

Bonsai.ML

Intelligent Experimental Control

Nicholas Guilbeault and Gonalo Lopes and Joaqu n Rapela

Gatsby Computational Neuroscience Unit
NeuroGEARS Ltd.

December 5, 2024



Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

History of Bonsai.ML

- ▶ grant application
- ▶ developing tools
- ▶ who cares about Bonsai.ML tools? → more focus on dissemination
- ▶ real-time ML

Goals of Bonsai.ML

- ▶ allow non-programmers use ML tools,
- ▶ learn from non-programmers what ML tools are useful for them

Need for Real-Time and Reactive ML


Conventional ML operates on stored datasets. We need real-time ML that operates on infinite data stream with time-varying statistical properties.

If a sensor fails, our inferences need to continue. That is, we need reactive ML (e.g., `rx.infer`).

Bonsai.ML demos

AEON project

We are building a new type of experimentation: continual recording of behavioral and neural data of mice foraging in large arenas for weeks to months.



BBSRC grant: Machine learning for Neuroscience experimental control

Abstract

To understand the brain, scientists aim to explain how animal behaviour relates to neural activity. This requires the design and precise control of behavioural experiments, wherein animals perform particular tasks while experimenters either record or manipulate neural activity in specific neural circuits. Such experiments require data acquisition software that integrates and controls hardware from multiple recording devices (cameras, electrodes, sensors), and analysis tools that can interpret large and complex datasets. Progress is held back by the lack of standardised tools for design and implementation of experimental protocols, and the difficulty of integrating state-of-the-art data processing and neuroinformatics into custom experimental designs. The fields of behavioural and brain sciences have consequently suffered from both inefficiency and poor reproducibility, due to disparate data acquisition and analysis solutions created independently across laboratories. To address these challenges, we propose to extend, enhance, maintain and support **Bonsai**, a fully integrated software environment to enable cutting-edge reproducible systems neuroscience experiments using animal models, with a particular emphasis on machine-intelligence-enabled, real-time neuroinformatics methods. While Bonsai is already adopted by hundreds of scientists worldwide, we aim to extend Bonsai's functionality with a toolbox of online and offline Machine Intelligence tools for analysis of behavioural and neural data (video-based analysis of behavioural motifs, real-time and offline analysis of neural signals), and create an open-access platform for software sharing and integration with multiple programming languages. Enhancing Bonsai's ecosystem will be a game-changer for behavioural and brain science experiments by enabling new types of research, increasing and diversifying user base, and dramatically improve efficiency and reproducibility of research.

[grant details](#)

Abstract

To understand the brain, scientists aim to explain how animal behaviour relates to neural activity. This requires the design and precise control of behavioural experiments, whereas achieving precise particular tasks while experiments either record or manipulate neural activity is equally neural activity. Such experiments require data acquisition software that integrates and controls hardware from multiple interacting devices (cameras, data stores, stimulus, and analysis) that can be deployed in a laboratory setting. This paper describes the design and implementation of Bonsai.ML, a software system for the design and implementation of experimental protocols, and the difficulty of integrating many of these data processing and manipulation tasks into a single experimental design. The field of Behavioural and brain sciences have consequently suffered from both inefficiency and poor reproducibility, due to disparate data acquisition and analysis systems created independently across laboratories. To address these challenges, we designed an experimental platform, Bonsai and support Bonsai.ML, a fully integrated software environment to create, manage, deploy, and execute experimental protocols using neural activity, with a particular emphasis on machine intelligence enabled, and data manipulation capability. Bonsai.ML is directly adaptable to hardware of scientific significance as per the experimental design. Bonsai.ML is a fully integrated software environment to create, manage, deploy, and execute experimental protocols using neural activity, with a particular emphasis on machine intelligence enabled, and data manipulation capability. Bonsai.ML is directly adaptable to hardware of scientific significance as per the experimental design.

grant details

- at the time of the grant I did not know much about Bonsai, and I did not want to know about it because it was a Windows software
- but the grant reviewers helped me understand how good Bonsai was :)
- controlling experiments using advanced machine learning methods
 - e.g., testing causality of neural patterns on behavior; deactivate a brain region when you forecast the appearance of a pattern of brain activity related to a given behavior, and check if after the deactivation the behavioral pattern disappears
 - e.g., only visually stimulate an animal until you achieve a certain precision in the receptive field estimation

My experience in machine learning and Neuroscience

I have developed methods for:

- ▶ nonlinear regression methods to estimate receptive fields of visual cells ([Rapela et al., 2006, 2010](#)),
- ▶ Bayesian linear regression methods to understand the relation between phase concentration in the human EEG and attention ([Rapela et al., 2012a,b, 2018](#)),
- ▶ dynamical systems models to model the relation between ECoG measurements and speech production in humans ([Rapela, 2016, 2017, 2018](#)),
- ▶ unsupervised models to characterize epilepsy using Utah array recordings from humans ([Rapela et al., 2019](#); [Rapela and Todorov, 2019](#)).

and I am the main developer of [svGPFA](#), a method using variational inference on Gaussian processes to infer latent variables from Neuropixels population recordings.

What type of machine learning we want for Bonsai?

supervised, unsupervised and reinforcement methods

supervised methods find mappings between inputs X and outputs y , like in curve fitting

unsupervised methods discover structure in inputs X , without any output, like in clustering

reinforcement learning methods learn relations between inputs X and actions a to maximize future rewards.

batch vs online processing

batch processing all data is read (generally from files) and processed at the same time.

online processing data is processed as it arrives and processed one at a time

What type of machine learning we want for Bonsai?

stationary vs non-stationary data

stationary data has statistics (i.e., characteristics) that do not change with time.

non-stationary data has time-varying statistics.

iid vs time-series datasets

iid datasets contain samples that are unrelated (i.e., independent) from each other and all come from the same distribution. For example a dataset of coin tosses is iid.

time-series datasets contain samples that are related to each other. For example a dataset of frames from a movie is a time-series one.

What type of machine learning we want for Bonsai?

probabilistic vs deterministic models

probabilistic models assume that variables of interest are random quantities, and seek to estimate their distribution.

deterministic models treat their variables of interest as deterministic quantities.

reactive vs non-reactive models

non-reactive inference models perform inference by following a pre-established sequence of steps, assuming availability of data before each step begins.

reactive inference models react to data availability performing a inference step only when data becomes available. Reactive inference models allow continuous inference in scenarios were data sources (e.g., cameras) can appear or disappear over time.

Machine learning models for Bonsai

For Bonsai we want ML models that:

can process online data process one data item at a time, when they are produced (reactively), and can handle infinite data streams

are non-stationary can process data with time-varying statistics

can handle time-series datasets as most neuroscience datasets are time series (e.g., behavioral videos, neuron spike counts).

are reactive can continue doing inference when adding/removing data sources

A new type of machine learning for Bonsai

Most machine learning methods used today:

- ▶ process **offline** (i.e., dead) finite data that has been saved to files,
- ▶ assume that the characteristics of data are constant across time (i.e., **stationarity assumption**).

A new type of machine learning for Bonsai

Most machine learning methods used today:

- ▶ process **offline** (i.e., dead) finite data that has been saved to files,
- ▶ assume that the characteristics of data are constant across time (i.e., **stationarity assumption**).

But Bonsai needs machine learning methods that:

- ▶ process **online** (i.e., live) potentially infinite datastreams,
- ▶ can operate on data with characteristics that change across time (i.e., **non-stationarity assumption**).

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion


Linear Regression: Fundamental Concepts

- ▶ main concepts of linear regression
- ▶ Online Bayesian Linear Regression (can process infinite data streams, but assumes stationarity)
- ▶ Recursive least squares (can process infinite data streams, and does not assume stationarity)

Linear Regression: Practical


Estimation of receptive fields of visual cells from the Allen Institute.

Linear regression example




Keshavarzi et al., 2021


Linear regression example




Keshavarzi et al., 2021



Linear regression example



Keshavarzi et al., 2021



Is there a linear relation between the speed of rotation and the firing rate of visual cells?


Estimating nonlinear receptive fields from natural images

$$y(\mathbf{x}) = k_0 + \sum_{i,j=1}^N k_1(i,j)x(i,j) \quad (1)$$

$$+ \sum_{i_1,j_1,i_2,j_2=1}^N k_2(i_1,j_1,i_2,j_2)x(i_1,j_1)x(i_2,j_2) + \dots$$

$$+ \sum_{i_1,j_1,\dots,i_Q,j_Q=1}^N k_Q(i_1,j_1,\dots,i_Q,j_Q)x(i_1,j_1)\dots x(i_Q,j_Q) + \varepsilon$$

$$y(\mathbf{x}) = \mathbf{A}\mathbf{q}(\mathbf{x}) + \varepsilon \quad (2)$$



Linear regression model

simple linear regression model

$$\begin{aligned}y(x_i, \mathbf{w}) &= w_0 + w_1 x_i = [1, x_i] \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = [\phi_0(x_i), \phi_1(x_i)] \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \\ &= \phi(x_i)^T \mathbf{w}\end{aligned}$$

polynomial regression model

$$\begin{aligned}y(x_i, \mathbf{w}) &= w_0 + w_1 x_i + w_2 x_i^2 + w_3 x_i^3 = [1, x_i, x_i^2, x_i^3] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} \\ &= [\phi_0(x_i), \phi_1(x_i), \phi_2(x_i), \phi_3(x_i)] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{bmatrix} = \phi(x_i)^T \mathbf{w}\end{aligned}$$

basis functions linear regression model

$$y(x_i, \mathbf{w}) = \phi(x_i)^T \mathbf{w} = \sum_{j=1}^M w_j \phi_j(x_i)$$

Linear regression model

$$\mathbf{y}(\mathbf{x}, \mathbf{w}) = \begin{bmatrix} y(x_1, \mathbf{w}) \\ y(x_2, \mathbf{w}) \\ \vdots \\ y(x_N, \mathbf{w}) \end{bmatrix} = \begin{bmatrix} \phi_1(x_1) & \phi_2(x_1) & \dots & \phi_M(x_1) \\ \phi_1(x_2) & \phi_2(x_2) & \dots & \phi_M(x_2) \\ \vdots & \vdots & \dots & \vdots \\ \phi_1(x_N) & \phi_2(x_N) & \dots & \phi_M(x_N) \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_M \end{bmatrix}$$
$$= \Phi \mathbf{w}$$

where $\mathbf{y}(\mathbf{x}, \mathbf{w}) \in \mathbb{R}^N$, $\Phi \in \mathbb{R}^{N \times M}$, $\mathbf{w} \in \mathbb{R}^M$.

Notes

- ▶ We learned how to build a linear regression model.
- ▶ But, how can we learn the model parameters w ?

Least-squares estimation of model parameters (Trefethen and Bau III, 1997)


Definition 1 (Least-squares problem)

Given $\Phi \in \mathbb{R}^{N \times M}$, $N \geq M$, $\mathbf{t} \in \mathbb{R}^N$, find $\mathbf{w} \in \mathbb{R}^M$ such that $E_{LS}(\mathbf{w}) = \|\mathbf{t} - \Phi\mathbf{w}\|_2$ is minimised.

Theorem 1 (Least-squares solution)

Let $\Phi \in \mathbb{R}^{N \times M}$ ($N \geq M$) and $\mathbf{t} \in \mathbb{R}^N$ be given. A vector $\mathbf{w} \in \mathbb{R}^M$ minimises $\|\mathbf{r}\|_2 = \|\mathbf{t} - \Phi\mathbf{w}\|_2$, thereby solving the least-squares problem, if and only if $\mathbf{r} \perp \text{range}(\Phi)$, that is, $\Phi^T \mathbf{r} = 0$, or equivalently, $\Phi^T \Phi \mathbf{w} = \Phi^T \mathbf{t}$, or again equivalently, $P\mathbf{t} = \Phi\mathbf{w}$, where $P \in \mathbb{R}^{N \times N}$ is the orthogonal projector onto $\text{range}(\Phi)$ (i.e., $P = A(A^T A)^{-1} A^T$).

Figure 3.2 Geometrical interpretation of the least-squares solution, in an N -dimensional space whose axes are the values of t_1, \dots, t_N . The least-squares regression function is obtained by finding the orthogonal projection of the data vector \mathbf{t} onto the subspace spanned by the basis functions $\phi_j(\mathbf{x})$ in which each basis function is viewed as a vector φ_j of length N with elements $\phi_j(\mathbf{x}_n)$.



Bonsai.ML

Linear Regression

Least-squares regression

Least-squares estimation of model parameters
(Trefethen and Bau III, 1997)

Least-squares estimation of model parameters (Trefethen and Bau III, 1997)

Definition 1 (Least-squares problem)

Given $\Phi \in \mathbb{R}^{N \times M}$, $N \geq M$, $\mathbf{t} \in \mathbb{R}^N$, find $\mathbf{w} \in \mathbb{R}^M$ such that $\|\Phi\mathbf{w} - \mathbf{t}\|_2$ is minimized.

Theorem 1 (Least-squares solution)

Let $\Phi \in \mathbb{R}^{N \times M}$, $N \geq M$, and $\mathbf{t} \in \mathbb{R}^N$ be given. A vector $\mathbf{w} \in \mathbb{R}^M$ minimizes $\|\Phi\mathbf{w} - \mathbf{t}\|_2 = \|\mathbf{t} - \Phi\mathbf{w}\|_2$ (minimizing the least-squares problem) if and only if $\mathbf{t} - \text{range}(\Phi)$, that is, $\Phi^T(\mathbf{t} - \Phi\mathbf{w}) = 0$, or equivalently, $\Phi^T\Phi\mathbf{w} = \Phi^T\mathbf{t}$, or again equivalently, $P\mathbf{t} = \Phi\mathbf{w}$, where $P \in \mathbb{R}^{N \times N}$ is the orthogonal projector onto $\text{range}(\Phi)$ (i.e., $P = A(A^TA)^{-1}A^T$).

Figure 1.2 Geometrical interpretation of the least-squares solution. In an N -dimensional space of data and an M -dimensional space of model parameters, the least-squares regression function is obtained by finding the orthogonal projection of the data vector \mathbf{t} onto the subspace spanned by the basis functions ϕ_1, \dots, ϕ_M in which each basis function is viewed as a vector ϕ_j of length 1 with elements $\phi_j(i)$.



Bishop (2011)

Given a set of N observations, \mathbf{t} , $N > M$, we want to find model parameters \mathbf{w} such that the model outputs, $\mathbf{t}(\mathbf{x}, \mathbf{w})$ equal the observations. This is generally impossible, because the degrees of freedom of the observations, N , is generally larger than the degrees of freedom of the model $\mathbf{t}(\mathbf{x}, \mathbf{w})$, M . We instead solve the following least-squares problem.

Instruction to run notebooks in Google Colab

1. open a notebook from [here](#)
2. replace **github.com** by **githubtocolab.com** in the URL
3. insert a cell at the beginning of the notebook with the following content

```
!git clone https://github.com/joacorapela/gcnuBridging2023.git
%cd gcnuBridging2023
!pip install -e .
```

4. from the menu **Runtime** select **Run all**.

Code for least-squares estimation of model parameters

- ▶ overfitting
- ▶ cross validation
- ▶ larger datasets allow more complex models

Notes

- ▶ We learned how to estimate the parameters of a linear regression models by least squares.
- ▶ But, how to avoid overfitting in the estimation?

Regularised least-squares estimation of model parameters

To cope with the overfitting of least squares, we can add to the least squares optimisation criterion a term that enforces coefficients to be zero. The regularised least-squares optimisation criterion becomes:

$$E_{RLS}(\mathbf{w}) = \|\mathbf{t} - \Phi\mathbf{w}\|_2^2 + \lambda\|\mathbf{w}\|_2^2$$

where λ is the regularisation parameter that weights the strength of the regularisation.

Regularised least-squares estimation of model parameters

Claim 1 (Regularised least-squares estimate)

$$\mathbf{w}_{RLS} = \arg \min_{\mathbf{w}} E_{RLS}(\mathbf{w}) = \arg \min_{\mathbf{w}} \|\mathbf{t} - \Phi \mathbf{w}\|_2^2 + \lambda \|\mathbf{w}\|_2^2 = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

Proof.

Since $E_{RLS}(\mathbf{w})$ is a polynomial of order two on the elements of \mathbf{w} (i.e., a quadratic form), we can use the *Completing the Squares* technique below to find its minimum.

$$\begin{aligned} \boldsymbol{\mu} &= \arg \max_{\mathbf{w}} \mathcal{N}(\mathbf{w} | \boldsymbol{\mu}, \Sigma) = \arg \max_{\mathbf{w}} \log \mathcal{N}(\mathbf{w} | \boldsymbol{\mu}, \Sigma) \\ &= \arg \max_{\mathbf{w}} \left\{ K - \frac{1}{2} (-2\boldsymbol{\mu}^T \Sigma^{-1} \mathbf{w} + \mathbf{w} \Sigma^{-1} \mathbf{w}) \right\} \end{aligned} \quad (3)$$

$$\begin{aligned} &= \arg \min_{\mathbf{w}} \left\{ -K + \frac{1}{2} (-2\boldsymbol{\mu}^T \Sigma^{-1} \mathbf{w} + \mathbf{w} \Sigma^{-1} \mathbf{w}) \right\} \\ &= \arg \min_{\mathbf{w}} \{ K_1 - 2\boldsymbol{\mu}^T \Sigma^{-1} \mathbf{w} + \mathbf{w} \Sigma^{-1} \mathbf{w} \} \end{aligned} \quad (4)$$

To find the minimum of a quadratic form, we write it in the form of the terms inside the curly brackets of Eq. 4, and the term corresponding to $\boldsymbol{\mu}$ will be the minimum.

Regularised least-squares estimation of model parameters

Proof.

Let's write E_{RLS} in the form of the terms inside the curly brackets of Eq. 4.

$$\begin{aligned} E_{RLS} &= ||\mathbf{t} - \Phi\mathbf{w}||_2^2 + \lambda ||\mathbf{w}||_2^2 = (\mathbf{t} - \Phi\mathbf{w})^T (\mathbf{t} - \Phi\mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w} \\ &= \mathbf{t}^T \mathbf{t} - 2\mathbf{t}^T \Phi\mathbf{w} + \mathbf{w}^T \Phi^T \Phi \mathbf{w} + \lambda \mathbf{w}^T \mathbf{w} \\ &= \mathbf{t}^T \mathbf{t} - 2\mathbf{t}^T \Phi\mathbf{w} + \mathbf{w}^T (\Phi^T \Phi + \lambda \mathbf{I}_M) \mathbf{w} \end{aligned}$$

Calling

$$\begin{aligned} \Sigma^{-1} &= \Phi^T \Phi + \lambda \mathbf{I}_M \\ \boldsymbol{\mu}^T \Sigma^{-1} &= \mathbf{t}^T \Phi \text{ or } \boldsymbol{\mu}^T = \mathbf{t}^T \Phi \Sigma \text{ or } \boldsymbol{\mu} = \Sigma \Phi^T \mathbf{t} = (\Phi^T \Phi + \lambda \mathbf{I}_M)^{-1} \Phi^T \mathbf{t} \end{aligned}$$

we can express

$$E_{RLS} = K + 2\boldsymbol{\mu}^T \Sigma^{-1} \mathbf{w} + \mathbf{w}^T \Sigma^{-1} \mathbf{w}$$

Then

$$\mathbf{w}_{RLS} = \arg \min_{\mathbf{w}} E_{RLS}(\mathbf{w}) = \boldsymbol{\mu} = (\Phi^T \Phi + \lambda \mathbf{I}_M)^{-1} \Phi^T \mathbf{t}$$

□

Code for regularised least-squares estimation of model parameters

- ▶ control of overfitting

Notes

- ▶ So far we have assumed deterministic parameters. But they are random quantities.
- ▶ How can we build linear regression models for random parameters?

Maximum-likelihood estimation of model parameters

Definition 2 (Likelihood function)

For a statistical model characterised by a probability density function $p(\mathbf{x}|\theta)$ (or probability mass function $P_\theta(X = \mathbf{x})$) the likelihood function is a function of the parameters θ , $\mathcal{L}(\theta) = p(\mathbf{x}|\theta)$ (or $\mathcal{L}(\theta) = P_\theta(\mathbf{x})$).

Definition 3 (Maximum likelihood parameters estimates)

The maximum likelihood parameters estimates are the parameters that maximise the likelihood function.

$$\theta_{ML} = \arg \max_{\theta} \mathcal{L}(\theta)$$

Maximum-likelihood estimation for the basis function linear regression model

We seek the parameter \mathbf{w}_{ML} and β_{ML} that maximised the following likelihood function

$$\mathcal{L}(\mathbf{w}, \beta) = p(\mathbf{t}|\mathbf{w}, \beta) = \mathcal{N}(\mathbf{t}|\Phi\mathbf{w}, \beta^{-1}I_N) \quad (5)$$

They are

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \quad (6)$$

$$\frac{1}{\beta_{ML}} = \frac{1}{N} \|\mathbf{t} - \Phi\mathbf{w}_{ML}\|_2^2 \quad (7)$$

- ▶ first regression method that assumes random observations
- ▶ if the likelihood function is assumed to be Normal, maximum-likelihood and least-squares coefficients estimates are equal.

Maximum likelihood: exercise

Exercise 1

Derive the formulas for the maximum likelihood estimates of the coefficients, \mathbf{w} , and noise precision, β , of the basis functions linear regression model given in Eqs. 6 and 7.

Solution.

$$\begin{aligned}\mathcal{L}(\mathbf{w}, \beta) &= p(\mathbf{t}|\mathbf{w}, \beta) = \mathcal{N}(\mathbf{t}|\Phi\mathbf{w}, \beta^{-1}\mathbf{I}) \\ &= \frac{1}{(2\pi)^{\frac{N}{2}} |\beta^{-1}\mathbf{I}|^{\frac{1}{2}}} \exp\left(-\frac{\beta}{2} \|\mathbf{t} - \Phi\mathbf{w}\|_2^2\right) \\ \log \mathcal{L}(\mathbf{w}, \beta) &= -\frac{N}{2} \log 2\pi + \frac{N}{2} \log \beta - \frac{\beta}{2} \|\mathbf{t} - \Phi\mathbf{w}\|_2^2 \\ \mathbf{w}_{ML} &= \arg \max_{\mathbf{w}} \log \mathcal{L}(\mathbf{w}, \beta) = \arg \min_{\mathbf{w}} \|\mathbf{t} - \Phi\mathbf{w}\|_2^2 = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \\ \frac{\partial}{\partial \beta} \log p(\mathbf{t}|\mathbf{w}_{ML}, \beta) &= \frac{N}{2} \frac{1}{\beta} - \frac{1}{2} \|\mathbf{t} - \Phi\mathbf{w}_{ML}\|_2^2 \\ \frac{\partial}{\partial \beta} \log p(\mathbf{t}|\mathbf{w}_{ML}, \beta_{ML}) &= 0 \quad \text{iff} \quad \frac{1}{\beta_{ML}} = \frac{1}{N} \|\mathbf{t} - \Phi\mathbf{w}_{ML}\|_2^2\end{aligned}$$

Notes

- ▶ We have learned how to estimate random parameters in linear regression models.
- ▶ How can we incorporate prior assumptions in this estimation?

Batch Bayesian linear regression: posterior distribution of parameters

In Bayesian linear regression we seek the posterior distribution of the weights of the linear regression model, \mathbf{w} , given the observations, which is proportional to the product of the likelihood function, $p(\mathbf{t}|\mathbf{w})$, and the prior, $p(\mathbf{w})$; i.e.,

$$p(\mathbf{w}|\mathbf{t}) \propto p(\mathbf{t}|\mathbf{w})p(\mathbf{w}) \quad (8)$$

To calculate this posterior below we use the likelihood function defined in Eq. 5 and the following prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|\mathbf{0}, \alpha^{-1}\mathbf{I})$$

Using the expression of the conditional of the Linear Gaussian model we obtain

$$p(\mathbf{w}|\mathbf{t}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_N, \mathbf{S}_N) \quad (9)$$

$$\mathbf{m}_N = \beta \mathbf{S}_N \Phi^T \mathbf{t}$$

$$\mathbf{S}_N^{-1} = \alpha \mathbf{I} + \beta \Phi^T \Phi \quad (10)$$

Batch Bayesian linear regression: exercise

Exercise 2

Derive the formulas for the Bayesian posterior mean (Eq. 9) and covariance (Eq. 10) of the basis function linear regression model.

Exercise 3

Show that

$$\log p(\mathbf{w}|\mathbf{t}) = K - \frac{\beta}{2} \|\mathbf{t} - \Phi\mathbf{w}\|_2^2 - \frac{\alpha}{2} \|\mathbf{w}\|_2^2 \quad (11)$$

Therefore, the maximum-a-posteriori parameters of the basis function linear regression model are the solution of the regularised least-squares problem with $\lambda = \alpha/\beta$.

Note that, as we will show next, Bayesian linear regression uses the full posterior of the parameters to make predictions or to do model selection, and not just the maximum-a-posteriori parameters.

Batch Bayesian linear regression: demo code

Available [here](#)

Notes

- ▶ Now we know how to do batch Bayesian linear regression.
- ▶ However, in Bonsai we don't want to work with batch data. We want to do online processing of infinite data streams. How can we do this?

Recursive update of posterior distribution of the parameters for conditionally independent observations

Claim 2 (recursive update)

If the observations, $\{\mathbf{t}_1, \dots, \mathbf{t}_n, \dots\}$, are linearly independent when conditioned on the model parameters, θ , then for any $n \in \mathbb{N}$

$$p(\theta|\mathbf{t}_1, \dots, \mathbf{t}_n) = K p(\mathbf{t}_n|\theta)p(\theta|\mathbf{t}_1, \dots, \mathbf{t}_{n-1}) \quad (12)$$

where K is a quantity that does not depend on θ .

Recursive update of posterior distribution of the parameters for conditionally independent observations

Proof.

By induction on $H_n : p(\theta | \mathbf{t}_1, \dots, \mathbf{t}_n) = K p(\mathbf{t}_n | \theta) p(\theta | \mathbf{t}_1, \dots, \mathbf{t}_{n-1})$.

H_1

$$p(\theta | \mathbf{t}_1) = \frac{p(\theta, \mathbf{t}_1)}{p(\mathbf{t}_1)} = \frac{p(\mathbf{t}_1 | \theta) p(\theta)}{p(\mathbf{t}_1)} = K p(\mathbf{t}_1 | \theta) p(\theta)$$

$H_n \rightarrow H_{n+1}$

$$\begin{aligned} p(\theta | \mathbf{t}_1, \dots, \mathbf{t}_{n+1}) &= \frac{p(\theta, \mathbf{t}_1, \dots, \mathbf{t}_{n+1})}{p(\mathbf{t}_1, \dots, \mathbf{t}_{n+1})} \\ &= \frac{p(\mathbf{t}_{n+1} | \theta, \mathbf{t}_1, \dots, \mathbf{t}_n) p(\theta, \mathbf{t}_1, \dots, \mathbf{t}_n)}{p(\mathbf{t}_1, \dots, \mathbf{t}_{n+1})} \\ &= \frac{p(\mathbf{t}_{n+1} | \theta) p(\theta, \mathbf{t}_1, \dots, \mathbf{t}_n)}{p(\mathbf{t}_1, \dots, \mathbf{t}_{n+1})} \\ &= \frac{p(\mathbf{t}_{n+1} | \theta) p(\theta | \mathbf{t}_1, \dots, \mathbf{t}_n) p(\mathbf{t}_1, \dots, \mathbf{t}_n)}{p(\mathbf{t}_1, \dots, \mathbf{t}_{n+1})} \\ &= K p(\mathbf{t}_{n+1} | \theta) p(\theta | \mathbf{t}_1, \dots, \mathbf{t}_n) \end{aligned}$$

Note: the third equality above holds because the observations are assumed to be conditional independent given the parameters.



Recursive update of the posterior distribution of the parameters for a conjugate prior

Claim 3

If

$$P(\mathbf{w}|\mathbf{t}_1, \dots, \mathbf{t}_n) = \mathcal{N}(\mathbf{w}|\mathbf{m}_n, \mathbf{S}_n) \quad (13)$$

$$P(\mathbf{t}_{n+1}|\mathbf{w}) = \mathcal{N}(\mathbf{t}_{n+1}|\Phi\mathbf{w}, \beta^{-1}\mathbf{I}) \quad (14)$$

then

$$P(\mathbf{w}|\mathbf{t}_1, \dots, \mathbf{t}_{n+1}) = \mathcal{N}(\mathbf{w}|\mathbf{m}_{n+1}, \mathbf{S}_{n+1})$$

with

$$\mathbf{S}_{n+1} = \mathbf{S}_n - (\beta^{-1} + \phi(\mathbf{x}_{n+1})^\top \mathbf{S}_n \phi(\mathbf{x}_{n+1}))^{-1} \mathbf{S}_n \phi(\mathbf{x}_{n+1}) \phi(\mathbf{x}_{n+1})^\top \mathbf{S}_n \quad (15)$$

$$\begin{aligned} \mathbf{m}_{n+1} = & \beta \mathbf{t}_{n+1} \mathbf{S}_{n+1} \phi(\mathbf{x}_{n+1}) + \mathbf{m}_n - \\ & (\beta^{-1} + \phi(\mathbf{x}_{n+1})^\top \mathbf{S}_n \phi(\mathbf{x}_{n+1}))^{-1} \phi(\mathbf{x}_{n+1})^\top \mathbf{m}_n \mathbf{S}_n \phi(\mathbf{x}_{n+1}) \end{aligned} \quad (16)$$


Python code for online Bayesian linear regression

Available [here](#).

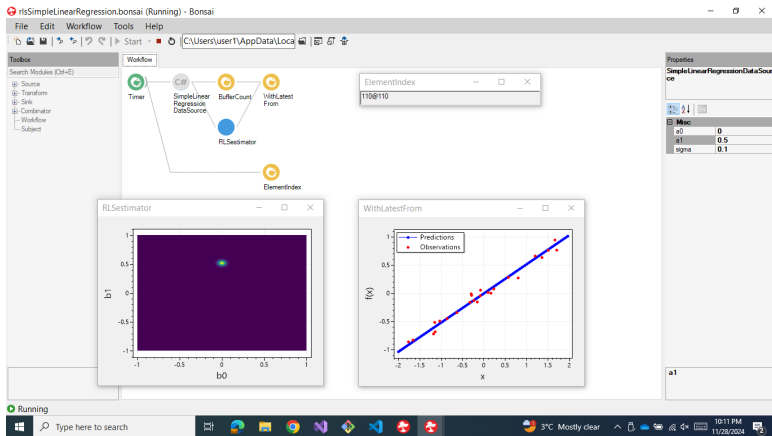
Notes

- ▶ We have learned how to do linear regression in Python.
- ▶ However, we are in the first Bonsai conference. Let's do online Bayesian linear regression in Bonsai.


Online Bayesian linear regression in Bonsai



Recursive least squares in Bonsai



Estimating receptive fields of cortical visual neurons in Bonsai



Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

Bonsai-Python Integration: Fundamental Concepts

- ▶ architecture
- ▶ async/sync regimes

Bonsai-Python Integration: Practical

- ▶ Bonsai-Python Hello World
- ▶ Bonsai-Python advanced example

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

Linear Dynamical Systems: Fundamental Concepts

Fundamental concepts of LDS.

Linear Dynamical Systems: Practical

Estimation of foraging mice kinematics.

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

Hidden Markov Models: Fundamental Concepts

Fundamental concepts of HMMs.

Hidden Markov Models: Practical

Estimation of behavioral states of foraging mice.

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

State-Space Decoders: Fundamental Concepts

Fundamental concepts of state-space decoders.

State-Space Decoders: Practical

Mice position decoding from hippocampal population activity.

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

Outline

Introduction

Linear Regression

Bonsai-Python Integration

Linear Dynamical Systems

Hidden Markov Models

State-Space Decoders

Roadmap

Discussion

Discussion

- ▶ how to disseminate Bonsai.ML?
 - ▶ distribute methods, with high-quality code, documentation and examples.
 - ▶ publish papers.
 - ▶ collaborate with experimentalists to use Bonsai.ML to tackle interesting neuroscience problems.
 - ▶ train the trainers.
 - ▶ train Bonsai users on basic machine learning topics.
 - ▶ find a killer Bonsai.ML application.
- ▶ suggestions for Bonsai.ML roadmap?
- ▶ suggestions for:
 - ▶ new ML models to integrate into Bonsai
 - ▶ new applications areas to investigate with Bonsai.ML
 - ▶ testing causality of brain activity patterns on behaviour
 - ▶ online selection of data to save (e.g. cameras)
 - ▶ Bonsai for online data analysis (e.g., Terry Sejnowski: for very large datasets, retrieval could be very onerous, so you'd better analyze data online).

References

- Bishop, C. M. (2016). *Pattern recognition and machine learning*. Springer-Verlag New York.
- Rapela, J. (2016). Entrainment of traveling waves to rhythmic motor acts.
- Rapela, J. (2017). Rhythmic production of consonant-vowel syllables synchronizes traveling waves in speech-processing brain regions.
- Rapela, J. (2018). Traveling waves appear and disappear in unison with produced speech.
- Rapela, J., Felsen, G., Touryan, J., Mendel, J. M., and Grzywacz, N. M. (2010). ePPR: a new strategy for the characterization of sensory cells from input/output data. *Network: Computation in Neural Systems*, 21(1-2):35–90.
- Rapela, J., Gramann, K., Westerfield, M., Townsend, J., and Makeig, S. (2012a). Brain oscillations in switching vs. focusing audio-visual attention. In *Proc. of 2012 IEEE Engineering in Medicine and Biology Society (EMBC'12)*, San Diego, CA.
- Rapela, J., Lin, T.-Y., Westerfield, M., Jung, T.-P., and Townsend, J. (2012b). Assisting autistic children with wireless EOG technology. In *Proc. of 2012 IEEE Engineering in Medicine and Biology Society (EMBC'12)*, San Diego, CA.
- Rapela, J., Mendel, J., and Grzywacz, N. (2006). Estimating nonlinear receptive fields from natural images. *Journal of Vision*, 6(4):441–474.
- Rapela, J., Proix, T., Todorov, D., and Truccolo, W. (2019). Uncovering low-dimensional structures of high-dimensional electrophysiological recordings in epilepsy. In *Engineering in Medicine and Biology Society (EMBC), 2019 41st Annual International Conference of the IEEE*, Berlin, Germany. IEEE.
- Rapela, J. and Todorov, D. (2019). Hidden markov models applied to LFPs from layer two and three of human cortex reveal highly stereotypical discrete states in epileptic seizures separated by more than an hour. Accepted at the 41st Engineering in Medicine and Biology Conference of the IEEE; main text https://sccn.ucsd.edu/~rapela/papers/EMBC19_Rapela_and_Todorov_EMBC_2019.pdf and supplemental information [embc19_hmm_supplemental.pdf](#).
- Rapela, J., Westerfield, M., and Townsend, J. (2018). A new single-trial foreperiod effect on inter-trial phase coherence. Part I: existence and relevance. *Neural Computation*, 30(9):2348–2383.
- Trefethen, L. n. and Bau III, D. (1997). Numerical linear algebra.